

Preston McCauley

Founder · Fractional AI Consultant · Product Strategist · Experience Design Expert

25+ Years

Strategy · Product · UX · AI

Top 75

Dallas Innovates AI Innovators

2 Books

Gen AI for Everyone · AI-First Mindset (writing)

65+ Reasoning Frameworks

CLEAR · CIC · SHAPE · QES + more



Preston McCauley is a Founder, Fractional AI Consultant, Product Strategist, and Experience Design Expert with over 25 years of experience spanning strategy, product, UX, and AI. As the founder of Clear Sight Designs, he specializes in building complex cognitive architectures and has authored two books, Gen AI for Everyone and AI-First Mindset. Recognized among Dallas Innovates' Top 75 AI Innovators, Preston has developed more than 65 reasoning frameworks—including CLEAR, CIC, SHAPE, and QES—that help organizations and individuals adopt an AI-first approach to solving meaningful problems.

Reach out I'd love to work with you!

 clearsightdesigns.com

 www.linkedin.com/in/preston-mccauley-immersive-ux/

 preston@clearsightdesigns.com


WHERE AGENTS BREAK

Context, Governance & Failure Surfaces in Agentic Workflows

Agents are everywhere now.

And somehow, they still aren't doing my laundry.

AI agents are already moving from pilots into production: Google Cloud's 2025 ROI of AI report says 52% of surveyed executives report deploying AI agents in production, while Gartner projects task-specific agents will be embedded in up to 40% of enterprise applications by 2026.



What's Your Biggest Agentic Fear?

What keeps you up at night?



Acts on Bad Context



Acts With Authority Beyond Parameters



Is Confident - Too Confident



Passes Bad Information

The Black Box Hasn't Disappeared it Moved



The workflow is now part of the black box.

Agents Are Already Doing the Work



The question is not:

Will agents do real work?

The question is:

Can we inspect what we are trusting?

In the US: roughly 45 million to 65 million active agentic systems now!

Approach With Care...


The agent gave me an answer. But...

Should I Trust It?

Was The Data Right?


Was Confidence Earned?

How Do I Know?



Two Kinds of Trust

Output Trust



Can I trust what it produced?

- Accurate
- Grounded
- Useful

Process Trust



Can I trust how it was created, routed, authorized, and passed forward?

- Clear
- Bounded
- Reviewed

Agentic systems need both.

Those Fears Have Names



Context Decay

Intent gets diluted across handoffs.



Authority Creep

Agents quietly expand what they are willing to do.



Trust Pressure

Fluency makes humans over-trust the output.



Downstream Propagation

The real danger appears where the output is trusted next.

The Dangerous Autonomous Moment

Ungoverned continuation

When the workflow continues simply because there is a next step.



Output moves forward



Assumptions move silently



Authority becomes implied



Risk travels downstream

Orchestration Is Not Governance



Orchestration says:

What happens next?

Route work · call tools · coordinate agents · manage state.



Governance asks:

Should it happen next?

- Under what authority?
- With what evidence?
- At what confidence?
- With what risk?
- With what audit trail?

ENTER THE CLEAR GOVERNANCE FRAMEWORK

What CLEAR Governance Is Actually Doing

CLEAR Governance makes the process inspectable. It asks:



Was the intent clear?



Were the limits enforced?



Was the evidence scored?



Was autonomy earned?



Was the result reviewed before moving downstream?

CLEAR as Runtime Questions



Clarity

Is the goal specific enough to act on?



Limits

What is allowed, blocked, or requires approval?



Examples / Evidence

What supports this step?



Adaptation


What autonomy has been earned?



Reflection

Is this safe to pass downstream?

Reasoning, Scoring, Execution



Reasoning can be agentic. **Scoring should be computational. Execution should be governed.**

One system reasoning, scoring itself, authorizing itself, and acting without separation isn't governance — it's generated confidence.

Pass Points: Receipts for Agent Handoffs



A pass point is a receipt for an agent handoff.

A governed handoff contract — output plus the trust state behind it.

It records:

What was completed

What was assumed

What evidence was used

What is allowed

What is blocked


What confidence exists

What autonomy is permitted


Whether human approval is required

Before / After

Typical agent workflow



Governed agent workflow



A normal workflow passes outputs. A governed workflow passes process state.

Did It Survive the Handoff?

Measuring governance drift across agent boundaries

Dimension	Before Agent		After Agent
Clarity	0.92	→	0.88
Limits	1.00	→	1.00
Evidence	0.74	→	0.82
Adaptation	0.64	→	0.64
Reflection	0.75	→	0.45 ⚠
Drift: 0.18 — Governance Warning			

Governance Drift: Why It Matters

Tracing shows what happened.

Drift measurement shows what degraded.

It can identify:



Which dimension changed



Which agent caused it



Whether review is required



Whether autonomy should be reduced



Whether the workflow should stop

From Output Trust to Process Trust



Output trust asks:

Is the result good?

Accurate · useful · grounded · relevant.



Process trust asks:

Did the result come from a governed path?

Evidence · risk · assumptions · authority to continue.

The strongest systems need both.

You Can Predict a Tornado — But Can You Prove It?

Tornado Onset flagged Union City, MI as CRITICAL hours before any official warning. The hard part isn't the call. It's proving the call existed before the storm.

MARCH 6, 2026



10:22 AM

Onset seals CRITICAL — Union City



1:45 PM

NWS issues first tornado watch



3:48 PM

Touchdown EF-3 · 160 mph

3 h 23 min before the NWS · 5 h 26 min before touchdown

The problem:

Claiming a prediction **after the fact is worthless** — hindsight is indistinguishable from foresight. Insurers won't stage adjusters on a claim; utilities won't move crews on a story. The prediction must be **provably sealed before the event**, in a tamper-proof chain anyone can independently verify. Change one field — a coordinate, a timestamp, a parameter — and the chain breaks. **The math is the proof.**

The Loop Makes The Gap Provable

By 10:22 AM the loop had completed all six turns and written the folder. The anecdote “we called it first” becomes a measurement.



Our seal — proven

- ✓ Sealed 10:22 AM (Genesis hash)
- ✓ Verified vs 4 SPC reports (Harvest hash)
- ✓ Chain INTACT · Grade B · conf 0.683
- ✓ CLEAR 1.00 · DRAFT, human-gated · 6 blocked



NWS — same moment



- ✗ Zero tornado watches
- ✗ Zero tornado warnings
- ✗ Zero severe warnings
- ✗ “Slight chance of severe storms”

5 h 26 min before touchdown · **3 h 23 min** before the first watch · **0.0 miles** from center. A measurement underwriters and emergency managers can build on.

CLEAR Governance Is a Loop That Drives the Agents

Not a scorecard at the end. The same five moves fire at every handoff — and decide whether the workflow earned the right to continue.

Think of it like a checkpoint the work clears at every handoff — not once at the door, but every single time it changes hands.



If not earned → BLOCKED

The pipeline stops. It doesn’t limp forward on weak governance.




Every turn leaves a trace

A file saved to disk — proof the decision was made, not just claimed.

Autonomy Is Routed, Not Assumed

How much leash does the AI get? Like a new hire — some calls they make alone, some they draft for your sign-off, some they're not allowed to touch. The score sets the leash.

`determine_autonomy()` picks the level from the score; weak governance is sent to BLOCKED.



Too weak? Stop / Clarify → **BLOCKED**. It halts instead of guessing forward.

Our storm run earned Level 4. It could prepare the sealed proof — but a human had to approve. A normal agent would've just published.

Six Turns of the Loop on One Prediction

Each turn does a governance job — in plain terms, then in code — and saves a trace.

1 Intent

Is the goal clear enough to act on?

evaluate_clear() → EXPLICIT 1.00. Vague would halt.

 CLEAR state • clarity score

2 Evidence

Find real proof — and you can't make it up.

Limits block fabrication. 4 SPC confirms → CLEAR 0.55→1.00.

 evidence_scores • clear_handoffs

3 Risk

Who gets hurt if this is wrong downstream?

assess_failure_surfaces() → downstream 0.70 (primary).

 failure-surface assessment

4 Authority

How much can it do, given the evidence?

12-input score + 3-model vote (2 of 3) → 0.683.

 governance_scorecard • judge_array

5 Route

Did it earn the next step, or stop?

0.683 + risk → Level 4 DRAFT. <0.40 would BLOCK.

 governance_decision

6 Seal

Lock it so it can't be faked later.

seal_genesis + seal_harvest → chain INTACT.

 sealed_packages • flow.log

What the Loop Leaves Behind

The output isn't a prediction. It's a folder of governed traces — each from a loop turn, each independently verifiable.




output/run_2026-03-06/

	clear_evaluations.json	Intent turn	clarity, classification, composite, allowed/blocked
	pass_points.json	Every turn	governed contracts at each handoff
	clear_handoffs.json	Evidence turn	before/after snapshots + drift score
	governance_scorecard.json	Authority turn	grade, confidence, autonomy, failure surfaces
	judge_array_result.json	Authority turn	3 model votes, quorum, consensus
	sealed_packages.json	Seal turn	Genesis + Harvest hashes, chain status
	weather_intelligence_report.md	Evidence turn	agent research — NWS state, SPC reports
	flow.log	All turns	every metric, stage, blocked action — audit trail

This folder is the product — governed proof the prediction was made correctly, verifiable without trusting the system that made it.

The Model Is Replaceable. The Governance Is Structural.

Swap the LLM — Qwen for Gemma, Gemma for GPT, GPT for whatever ships next quarter. The governance does not move.



The model can change. The governance state cannot disappear.

Would This Apply to My System?

Anywhere agent output becomes someone else's input:



Code review



DevOps remediation



Support escalation



Legal drafting



Healthcare intake



Financial analysis



Enterprise research



Sales operations



Security triage



Internal knowledge
workflows

The Pattern You Can Use Monday

- 1 Find one agent handoff
- 2 Add a pass point
- 3 Capture intent, evidence, limits, risk, confidence, and approval state
- 4 Measure drift before and after the agent boundary
- 5 Route autonomy from that state
- 6 Log or seal the handoff

Govern the process before you trust the output.

Orchestration says what happens next.

Governance says whether it should happen next.